

# Comparative Analysis of Image Classification Algorithms on the Fashion-MNIST Dataset

Shuvam Chowdhury  
Machine Learning & Data Science  
Email: kshuvam63@gmail.com  
GitHub: <https://github.com/TheShuvam>

**Abstract**—The Fashion-MNIST dataset is used to train and test various image classification models and compare their performance. The algorithms discussed in this report are K-Nearest Neighbors (KNN), Logistic Regression, Support Vector Machines (SVM), Convolutional Neural Networks (CNN), and the zero-shot CLIP model. Basic exploratory data analysis (EDA) addresses class imbalances, differing pixel intensities, normalization, and dimensionality reduction. Each algorithm is evaluated using accuracy, precision, recall, and F1-score. Results show that CNN best performs image classification on the Fashion-MNIST dataset with metrics of 92.07%, 0.9208, 0.9207, and 0.9197 respectively. This paper provides reasoning for algorithm selection, comparison, and reproducible experimentation.

**Index Terms**—Fashion-MNIST, image classification, k-Nearest Neighbors, Logistic Regression, Support Vector Machine, Convolutional Neural Network, CLIP, zero-shot learning, transformer models

## I. INTRODUCTION

Fashion-MNIST is a dataset of 28x28 grayscale article images each with an associated label from one of 10 classes: tshirt/top, trouser, pullover, dress, coat, sandal, shirt, sneaker, bag, ankle boot. The dataset takes the same format as the MNIST dataset with 60,000 examples in the training set, and 10,000 examples in the test set. In this project, the Fashion-MNIST is further divided to include a validation set. The validation set pulls 10,000 samples from the training set, and the training set becomes 50,000 samples. An example of 25 random images from the training set is shown in Figure 1.

In this project, image classification is applied to the FashionMNIST dataset with both baseline and advanced models to ultimately compare model performance. In order, this paper reviews the machine learning algorithms K-Nearest Neighbors (KNN), Logistic Regression, and Support Vector Machines (SVM) along with the deep learning Convolutional Neural Network (CNN) algorithm. In addition to these supervised learning models, we also evaluate CLIP, a modern transformer-based vision–language model, in a zero-shot classification setting for comparison. The implementation for this project is available in the project’s GitHub repository [here](#).

## II. LITERATURE REVIEW

Image classification aims to identify features in an image and accurately assign that image to a class defined by one or



Fig. 1. 25 Random images selected from the Fashion-MNIST training Dataset

more characteristics. Recent advancements in vision–language learning have introduced models capable of zero-shot image classification without supervised training. One of the most notable architectures is CLIP (Contrastive Language–Image Pre-training), which learns a joint embedding space between images and natural-language descriptions. CLIP has demonstrated strong generalization across diverse datasets by leveraging large-scale pre-training on image–text pairs, enabling it to recognize visual concepts from textual prompts rather than task-specific labels [10]. This motivates its inclusion in this study as a modern baseline for comparison against traditional machine learning models and CNNs.

Data used in this procedure undergoes unique pre-processing which involves segmenting each image by the

normalization of its pixels. Unfortunately, flawed pictures such as those that are fragmented, noisy, or ambiguous, are inevitable. CNNs are said to manage the above factors well so are widely used in image classification.

CNN extracts characteristics from images to obtain information. Therefore, features are not passed to CNN but are discovered while the network trains on a set of images. In an experiment on the CIFAR-10 dataset, when compared against VGG16, another classifier that utilizes 3x3 convolution layers and a 2x2 max pool layer, CNN resulted in a 92.22VGG16 only performed with an 89.03scaledependent pooling method of CNN improves its accuracy, but consequently also increases computational time. Computational time can be decreased with advanced architecture left undiscussed.

CNN is applied to the Fashion-MNIST dataset in our report for image classification and is discovered to perform best against other tested algorithms. The described experiment supports our conclusion that CNN is the best fit for classifying our set of images as described further along in this paper.

In a published application of CNN, the neural network is applied to Korean food images taken from an AI-Hub platform in an attempt to expand the knowledge of Korean food image classification [5]. To begin, the images underwent the aforementioned unique pre-processing for image classification to improve accuracy. The data pre-processing that will be described for use on the Korean food images dataset is like the pre-processing we executed on the Fashion-MNIST dataset.

The Korean food image dataset was divided into a training, testing, and validation set with 105,427, 30,122, and 15,061 images respectively. The images were resized to 331x331 pixels each. Then, the newly pre-processed image set was augmented to counteract overfitting. Performing data augmentation produced images with varying levels of brightness, contrast, saturation, and alignment. Features were extracted from augmented image datasets and pre-trained with deep neural networks like ResNet- 50, RestNet-101, RestNet-152, MobileNetV2, InceptionResNetV2, and NasNetLarge [5].

The pre-trained deep neural networks mentioned above were connected to a CNN used to classify objects of the images in this project. Each classification model passed through an initial training stage and a fine-tuning stage. Then, the model was evaluated under the accuracy measurement  $TP + TN / TP + FP + TN + FN$  where TP is true positive, TN is true negative, FP is false positive, and FN is false negative. The study concluded that accuracy was consistent despite the network choice for feature extraction. The classification model repeatedly incorrectly distinguished images that combined food classes or included food with similar shapes and colors [5]. The complexity of each Korean food image makes classification difficult. This study does not resolve the consistent misclassification of some Korean food classes, so could be built upon in future work.

The Fashion-MNIST dataset deals with much less complex images than the Korean food image dataset, but understanding the process of classifying a larger and more diversified dataset was useful when applying the same technique on our simpler

scale.

### III. METHODOLOGY

#### A. Data Pre-processing

I begin by importing the Fashion-MNIST dataset into a Jupyter notebook and dividing it into training, validation, and testing sets containing 60,000 and 10,000 images respectively. From the training set, I separate 10,000 samples to form a validation set, leaving 50,000 samples for training. Each grayscale image of size  $28 \times 28$  pixels is flattened into a 784-dimensional feature vector. These flattened vectors are stored in Pandas DataFrames for the training, validation, and test sets, with an additional column, label, storing the true class for each image.

To facilitate model training, I normalize pixel values using the MinMaxScaler from scikit-learn, which scales all pixel intensities to the range  $[0, 1]$ . This normalization is appropriate because Fashion-MNIST images contain bounded grayscale values, ensuring no outliers and stable optimization behavior for models sensitive to feature magnitude (e.g., logistic regression, KNN, and SVM).

#### B. Exploratory Data Analysis

I conduct exploratory data analysis (EDA) to understand dataset characteristics and ensure that preprocessing and modeling choices are appropriate. Class distributions are computed for the training, validation, and test sets to verify dataset balance. All three partitions contain nearly identical class proportions, meaning no additional resampling or class-weighting strategies are required.

Next, I examine pixel-intensity distributions across the dataset. Since Fashion-MNIST consists of grayscale images, the pixel values reflect brightness and contrast differences between classes. These distributions are moderately skewed, justifying the use of range normalization. I also visualize example images from each class to verify that the dataset contains clear and distinguishable patterns.

Principal Component Analysis (PCA) is then applied to the normalized training data to reduce dimensionality and mitigate the curse of dimensionality in traditional machine-learning models. I compute the cumulative explained variance curve and determine that 84 principal components capture approximately 90% of the variance. The data is transformed to retain only these 84 components. These 84 components are used as features for logistic regression, KNN, and SVM. I do not use PCA for the convolutional neural network (CNN), which learns spatially coherent features directly from raw pixel grids. CNNs are designed to learn their own optimal feature representations and hierarchical features from raw pixel data. PCA does not account for the spatial relationships present in image data that CNN relies on for image classification.

#### C. Model Training

After performing normalization and principal component analysis (PCA), I train several supervised learning models on the Fashion-MNIST dataset: logistic regression,  $k$ -nearest

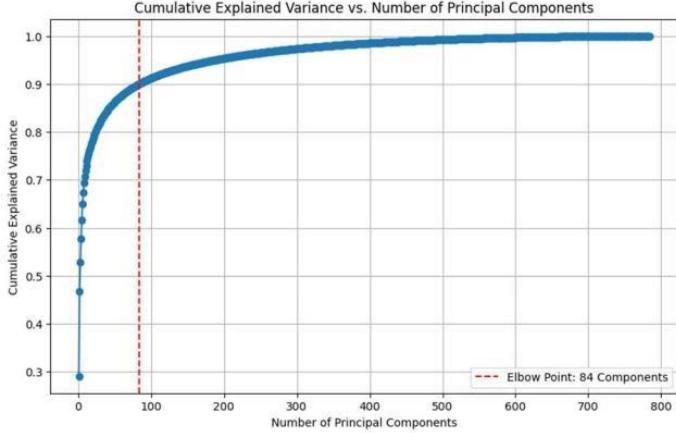


Fig. 2. Cumulative explained variance as a function of the number of PCA components. The dashed line indicates the elbow point at 84 components.

neighbors (KNN), support vector machines (SVM), and a convolutional neural network (CNN). The PCA-reduced features (84 components) are used for the logistic regression, KNN, and SVM models, while the CNN operates directly on the normalized image grids.

1) *Logistic Regression*: I first train a multinomial logistic regression classifier as a linear baseline. The model is fit on the PCA-reduced training data using the scikit-learn implementation with a maximum of 1000 iterations to ensure convergence. Let  $\mathbf{z} \in \mathbb{R}^{84}$  denote the PCA feature vector and  $y \in \{0, \dots, 9\}$  the class label. The model estimates class probabilities via the softmax function,

$$p(y = c \mid \mathbf{z}) = \frac{\exp(\mathbf{w}_c^\top \mathbf{z} + b_c)}{\sum_{j=0}^9 \exp(\mathbf{w}_j^\top \mathbf{z} + b_j)},$$

and predicts the class with maximum probability. I evaluate the trained model on the validation and test sets using accuracy, precision, recall, and F1-score.

2) *k-Nearest Neighbors (KNN)*: For the KNN classifier, I perform hyperparameter tuning over different values of  $k$  to identify the number of neighbors that yields the highest validation accuracy. Using Euclidean distance in the PCA feature space, each test sample is assigned the most frequent label among its  $k$  nearest neighbors in the training set. I sweep  $k$  over a predefined range, record the validation accuracy for each value, and select the  $k$  that maximizes validation performance. The final KNN model with the best  $k$  is then evaluated on the held-out test set.

3) *Support Vector Machine (SVM)*: I next train support vector machine classifiers using the PCA-reduced features. A grid search is conducted over kernel types (linear, polynomial, and radial basis function) and regularization strengths  $C$ . For each (kernel,  $C$ ) pair, I fit an SVM on the training data and compute validation accuracy. The combination that achieves the best validation performance is chosen as the final model. The selected SVM is then evaluated on the test set, and decision scores are used to compute receiver operating

characteristic (ROC) curves and area-under-the-curve (AUC) values in a one-vs-rest multi-class setting.

4) *Convolutional Neural Network (CNN)*: To exploit spatial structure in the images, I train a custom convolutional neural network directly on the normalized  $28 \times 28$  pixel grids. The input images are reshaped to tensors of size  $1 \times 28 \times 28$  and batched using PyTorch data loaders. The CNN consists of stacked convolutional layers with ReLU activation and max-pooling, followed by fully connected layers and a final softmax output over the ten classes. The network parameters are optimized using stochastic gradient descent (or Adam) with cross-entropy loss. I monitor training and validation accuracy over epochs to detect overfitting and select the model checkpoint with the best validation performance. Finally, I evaluate the CNN on the test set and compute accuracy, precision, recall, F1-score, and the confusion matrix for detailed error analysis.

#### D. CLIP Zero-Shot Classification

I evaluate three CLIP model variants for zero-shot classification on the Fashion-MNIST dataset. Among the tested models, the `openai/clip-vit-base-patch16` variant achieves the highest performance with a top-1 accuracy of 0.6803. After selecting this variant, I further refine the text prompts associated with each class. The prompt “a centered image of a {class}” yields the best results, improving the overall accuracy to 0.7115.

Despite these improvements, CLIP’s zero-shot performance remains significantly lower than that of the supervised models evaluated in this paper. This is expected, as CLIP is not fine-tuned on Fashion-MNIST and must rely solely on its large-scale pre-trained representations, which are optimized for natural images rather than grayscale clothing silhouettes.

The confusion matrix illustrates that CLIP struggles most with distinguishing between visually similar categories, such as shirts, T-shirts, and pullovers. This aligns with the model’s lack of exposure to Fashion-MNIST-specific visual patterns. While CLIP demonstrates reasonable generalization, its performance confirms that zero-shot vision–language models are not directly competitive with supervised approaches on this dataset.

## IV. EXPERIMENTAL RESULTS

The accuracy score referred to when evaluating the following algorithms is imported from `sklearn.metrics` and is a measure of subset accuracy where the set of labels predicted for a sample must exactly match the corresponding set of true labels [1].

### A. K-Nearest Neighbors

The goal of first employing the KNN algorithm is to achieve reasonable accuracy to justify more complex models if they outperform KNN significantly. The KNN algorithm deals with non-parametric learning. The value of  $k$ , where  $k$  defines the number of nearest training samples’ class labels to a test sample, is the only hyperparameter. The Fashion-MNIST dataset does not require extensive pre-processing, and

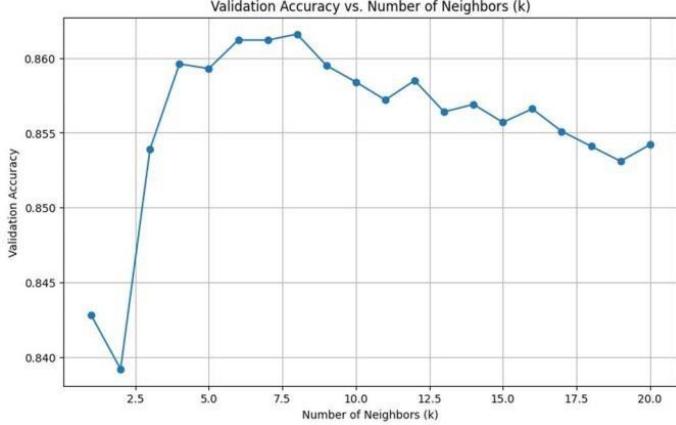


Fig. 3. Validation accuracy as a function of the number of neighbors ( $k$ ) in the K-Nearest Neighbors (KNN) classifier. The plot shows the performance trend across varying  $k$  values and highlights the range where the classifier achieves its highest accuracy.

KNN works well even with basic normalization. Features like pixel intensities or PCA-transformed data can be effectively compared using distance metrics.

To tune the hyperparameter, the training data is fit to multiple KNN classifiers with varying values of  $k$  neighbors and the validation accuracy vs. number of neighbors is plotted. The maximum point on the plot, (8, 0.86), represents the optimal number of neighbors and the highest achieved accuracy.

The KNN classifier is initialized and set with 8 optimal numbers of neighbors. A Stratified K-Fold Cross Validation is defined to achieve a more reliable estimate of model performance and preserve the proportion of each class in the training and validation sets by using five splits of the data. The KNN model is fit on the training set and predictions are made on the test set. The metrics calculated for the cross-validated sets equal the metrics returned when calculated on the test set. Accuracy, precision, recall, and F1-score are all 0.86. Therefore, the model is consistent across different data splits and correctly predicts the image class 86% of the time.

Take note that KNN has the most difficult time predicting class 6 (shirt) which it most confuses with class 0 (t-shirt/top).

#### B. Logistic Regression

Fashion-MNIST is a relatively low-complexity dataset. Where more complex machine learning models are likely to overfit, logistic regression is a simple model that does not require higher computational power and is more capable of generalizing. Logistic regression works well for linearly separable datasets, so the goal of testing this algorithm is to check how Fashion-MNIST performs under assumptions of linearity.

The logistic regression model is fit to the training data. The trained model is used to make predictions on the validation set first. Then, the model makes predictions on the test set and weighted averaging is used to compute accuracy,

Confusion Matrix on KNN Classifier										
Actual	0	1	2	3	4	5	6	7	8	9
0	851	0	12	18	6	1	99	1	12	0
1	6	966	3	15	4	0	5	0	1	0
2	21	1	775	9	115	0	74	0	5	0
3	33	4	14	874	42	0	30	0	3	0
4	1	0	114	30	779	0	72	0	4	0
5	1	0	0	0	0	896	0	59	2	42
6	181	2	123	21	88	0	570	0	15	0
7	0	0	0	0	0	9	0	955	0	36
8	0	1	8	4	7	0	10	6	962	1
9	0	1	2	3	4	5	6	7	8	9

Fig. 4. Confusion matrix of the K-Nearest Neighbors (KNN) classifier on the Fashion-MNIST test set. Diagonal entries indicate correct classifications, while off-diagonal values show misclassifications across the ten clothing categories.

precision, recall, and F1 scores to handle class imbalance. The pseudocode and its output for the outlined process are below.

```
log_reg = LogisticRegression(max_iter=1000,
                             random_state=0)
log_reg.fit(X_train, y_train)

y_val_pred = log_reg.predict(X_val)
# Compute classification metrics for the validation
# set

y_test_pred = log_reg.predict(X_test)
# Compute classification metrics for the test set

Output:
Validation Accuracy for Logistic Regression: 0.84
Validation Precision for Logistic Regression: 0.84
Validation Recall for Logistic Regression: 0.84
Validation F1 Score for Logistic Regression: 0.84
```

The logistic regression model performs slightly worse than the KNN model with only 84% accuracy. From the confusion matrix, the model appears to have an even more difficult time identifying class 6 (shirt) than KNN. In general, this may imply the data is not well suited to linear classification.

#### C. Support Vector Machine

SVM is effective for small to medium-sized datasets in high-dimensional spaces like Fashion-MNIST. SVM employs the kernel trick to transform the given data and find optimal decision boundaries. The kernel does not have to be linear, so SVM can capture more complex relationships between data points. SVM is utilized to compare Fashion-MNIST performance with no linear assumptions to the output of logistic regression which assumes linearity.

For hyperparameter tuning, a for loop iterates over all the combinations of two selected kernel types, radial basis function, and polynomial, and selected C values of 0.1, 1, and




Fig. 5. Confusion matrix of the Logistic Regression classifier evaluated on the Fashion-MNIST test set. Diagonal values represent correct predictions across the ten clothing categories, while off-diagonal entries indicate misclassifications.

10. The model is trained on the training set with each unique hyperparameter combination then accuracy is calculated by using the validation set for model prediction. The kernel, C value, and accuracy are stored in the empty list, result. Then, the highest accuracy is pulled from the results list, and the optimal kernel and C values are extracted. The pseudocode for this process is as follows.

```

result = []
kernel_options = [...]
C_values = [...]

for kernel in kernel_options:
    for C in C_values:
        svc_clf = SVC(kernel=kernel, C=C)
        svc_clf.fit(X_train, y_train)

        y_val_pred = svc_clf.predict(X_val)
        accuracy = accuracy_score(y_val, y_val_pred)

        # Append (kernel, C, accuracy) to result list
        result.append((kernel, C, accuracy))

# Identify the entry with the highest validation accuracy
best_result = max(result, key=lambda x: x[2])
best_kernel = best_result[0]
best_C = best_result[1]

```

The determined hyperparameters are used to initialize the SVM classifier, train the model on training data, create model predictions on testing data, and output 90accuracy. The nonlinear decision boundaries achieved by SVM are better suited to the Fashion-MNIST dataset than the linear separation performed by logistic regression.

The ROC curve illustrates the model's ability to distinguish between positive and negative classes. It measures the ranking of probabilities for the true class. Therefore, even if the model makes an incorrect prediction, but only assigns the true class a




Fig. 6. ROC curves for the multi-class classification task on the Fashion-MNIST dataset. Each curve represents a single class using a one-vs-rest approach, with corresponding AUC values shown in the legend. The diagonal dashed line indicates random performance, while the dotted line shows the ideal classifier.

slightly lower probability, it can achieve a high area under the curve (AUC). This explains the discrepancy between accuracy and AUC. If the AUC is equal to 1, this represents a perfect fit. This plot shows the performance of the SVM classifier across each class in the testing set. Specifically, it evaluates how well the classifier distinguishes each class from the others using a one-vs-all approach. Notice all the classes have an AUC extremely close or equal to one.

#### D. Convolutional Neural Network

The architecture of CNN with an input and output layer and many hidden layers provides high accuracy for image classification. Each layer learns to identify different features, and because all hidden neurons in the same layer share weights and bias values, the network is tolerant to object translation within an image. CNN does not require manual feature extraction, so the data pre-processing step may be simplified for cases where CNN is used. CNN is tested to determine the accuracy of an image classification method that requires less feature engineering.

The CNN algorithm is tested with 1, 2, 3, and 4 hidden layers and begins with the raw flattened images derived from Fashion-MNIST. Each pixel is normalized and takes a value between 0 and 1. The objective function used for CNN is crossentropy loss, and the Adam optimizer is used to update the model's parameters based on the gradients computed during backpropagation.

For each CNN algorithm tested on the Fashion-MNIST dataset, the model is trained in a loop for 10 epochs. After each epoch, the accuracy and loss measurements are printed. After the training loop, the model is set to evaluation mode where the loss on the validation set is calculated. Finally, the




Fig. 7. Confusion matrix of the 3-layer Convolutional Neural Network (CNN) classifier evaluated on the Fashion-MNIST test set. Diagonal entries represent correct predictions, while off-diagonal values indicate misclassifications across the ten clothing categories.

model is assessed on test data, and the classification metrics are output. For CNN with a single layer, evaluation mode on the validation set following training is excluded. When evaluated on the test data, the CNN algorithm with 3 hidden layers performs best with accuracy, precision, recall, and F1 scores of 92.07%, 0.9208, 0.9207, and 0.9197 respectively. Since all the classification metrics for this model exceed those of the algorithms explained above, it is concluded that CNN most accurately performs image classification on the FashionMNIST dataset. This can be further confirmed by examining the confusion matrix.

This confirms CNN superiority over classical ML models.

#### E. Contrastive Language-Image Pretraining

Contrastive Language-Image Pretraining (CLIP) was applied to Fashion-MNIST to perform zero-shot classification. Zero-shot classification refers to the machine learning technique where a model can classify data without being trained on those classes. As a multimodal model, CLIP is capable of processing given images paired with natural language descriptions to label the data.

Three CLIP model variants were tested. Model variant “openai/clip-vit-base-patch16” performed best with 0.6803 accuracy. After choosing the proper variant, prompt testing increased the model accuracy to 0.7115 when the prompt “a centered image of a ” was used. However, even with a proper prompt, the CLIP model did not produce a competitive accuracy score. The confusion matrix illustrates the model has the most difficulty classifying a shirt. This class has the most incorrect predictions across all methods described in this paper.

#### V. BUSINESS INSIGHTS

For retailers, being able to analyze images of clothing can enhance their business capabilities. With all its potential,




Fig. 8. Confusion matrix of the CLIP model evaluated on the Fashion-MNIST test set using the best-performing prompt. The matrix shows per-class prediction performance across all ten categories, with diagonal values indicating correct predictions and off-diagonal entries showing misclassifications.

image classification can aid in marketing, recommending merchandise to online shoppers through a visual search, keeping up with the latest fashion trends, providing insights for supply and demand, and generally helping improve business functionality [6].

Although CLIP provides an interesting zero-shot alternative and performs well on natural-image benchmarks, its accuracy on the grayscale Fashion-MNIST dataset is significantly lower than supervised approaches. Because CLIP is not fine-tuned for this domain, it is not a suitable candidate for operational deployment in business settings where reliability and consistency are essential. The CNN model remains the most practical and effective choice for real-world applications.

#### VI. CONCLUSION

In summary, the convolutional neural network achieved the highest performance on the Fashion-MNIST dataset, outperforming all classical machine learning methods considered in this study. Logistic regression, KNN, and SVM provided strong baselines when combined with PCA dimensionality reduction, but none matched the feature-learning capabilities of the CNN. The zero-shot CLIP model, while demonstrating promising generalization in natural-image domains, performed substantially worse due to its lack of fine-tuning and the limited semantic richness of grayscale Fashion-MNIST images. Including CLIP in this comparison highlights the contrast between traditional supervised learning and modern vision-language models, offering insight into where each paradigm succeeds and where domain-specific training remains essential.

#### REFERENCES

- [1] “accuracy\_score,” Scikit-learn. Available: [https://scikit-learn.org/dev/modules/generated/sklearn.metrics.accuracy\\_score.html](https://scikit-learn.org/dev/modules/generated/sklearn.metrics.accuracy_score.html) (accessed Nov. 19, 2024).

- [2] A. Fulare, K. Raghavendra Prasad, P. Johri, I. S. Abdulrahman, K. B. D. Arasi, and G. Guwalani, "Convolutional Neural Network Based Image Classification and Its Comparison with VGG-16 to Measure Accuracy," in *Proc. 2023 3rd Int. Conf. Advance Computing and Innovative Technologies in Engineering (ICACITE)*, Greater Noida, India, 2023, pp. 367–370. doi: [10.1109/ICACITE57410.2023.10183128](https://doi.org/10.1109/ICACITE57410.2023.10183128).
- [3] "Fashion-MNIST," GitHub. Available: <https://github.com/zalandoresearch/fashionmnist> (accessed Oct. 23, 2024).
- [4] G. L., "Why use SVM?," Alteryx Community. Available: <https://community.alteryx.com/?category.id=external> (accessed Nov. 18, 2024).
- [5] M. Young, *The Technical Writer's Handbook*. Mill Valley, CA: University Science, 1989.
- [6] M. Chun, H. Jeong, H. Lee, T. Yoo, and H. Jung, "Development of Korean Food Image Classification Model Using Public Food Image Dataset and Deep Learning Methods," *IEEE Access*, vol. 10, pp. 128732–128741, 2022. doi: [10.1109/ACCESS.2022.3227796](https://doi.org/10.1109/ACCESS.2022.3227796).
- [7] S. Kholin and A. Bitkina, "Image Classification: 6 Industries & 26 Use Cases You Can Try," Onix. Available: <https://onixsystems.com/blog/imageclassification-use-cases> (accessed Nov. 20, 2024).
- [8] "What Is a Convolutional Neural Network?," MathWorks. Available: <https://www.mathworks.com/discovery/convolutional-neural-network.html> (accessed Nov. 19, 2024).
- [9] "Why does Logistic Regression perform better than machine learning models in clinical prediction studies?," StackExchange. Available: <https://datascience.stackexchange.com/questions/120262/why-does-logistic-regression-perform-better-than-machine-learning-models-in-clin> (accessed Nov. 18, 2024).
- [10] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark, G. Krueger, and I. Sutskever, "Learning Transferable Visual Models From Natural Language Supervision," in *Proc. 38th Int. Conf. Machine Learning (ICML)*, 2021, pp. 8748–8763. Available: <https://arxiv.org/abs/2103.00020>.